

Querying Data on the Web

Alvaro A A Fernandes
School of Computer Science
University of Manchester

COMP62421:: 2015-2016

Part I

The Relational/SQL Setting [1]



Teaching Day 1 Outline



- 1 Introduction
- 2 Database Management Systems [1]
 - DBMS: Definition
 - DBMS: Languages
 - DBMS-Centred Applications
- 3 Database Management Systems [2]
 - DBMS: Internal Architecture
 - DBMS: Strengths, Weaknesses, Trends
 - DBMS: Architectural Variations
- 4 The Relational Case
 - The Relational Model
 - Relational Databases
- 5 Relational Query Languages [1]
 - Query Languages: Special Features
 - Relational Languages Compared, and Views
 - The Tuple Relational Calculus
 - A Relational Algebra [1]
- 6 Relational Query Languages [2]
 - A Relational Algebra [2]
 - RA Examples
 - SQL
- 7 Conclusion

This Week: Topics



- We revise the basic notions of
 - what a **database management system** (DBMS) is
 - the role that is played by the languages a DBMS supports
 - how DBMS-centred applications are designed
- We adopt an approach to describing the internal architecture of DBMSs that allows us to discuss
 - the strengths and weaknesses of classical DBMSs
 - the recent trends in the way organizations use DBMS
 - how architectures have been diversifying recently
- We revise the various kinds of query languages by looking into
 - the (tuple) relational calculus
 - a relational algebra
 - SQL

This Week: Learning Objectives



- We aim to revise and reinforce undergraduate-level understanding of DBMS as software systems, rather than as software development components.
- We aim to get under way in our postgraduate-level exploration of query languages, in preparation for further delving into query processing techniques next week.

Lecture 1

Database Management Systems [1]



Section 1

DBMS: Definition



Database Management Systems

What are they?

Definition

A **database management system (DBMS)** is

- a software package
- for supporting applications
- aimed at managing very large volumes of data
- efficiently and reliably
- transparently with respect to the underlying hardware and network infrastructures
- and projecting a coherent, abstract model
- of the underlying reality of concern to applications
- through high-level linguistic abstractions.

Managing Very Large Volumes of Data

What does it mean?

- managing** storing for querying (often, updating as well, and more recently searching, exploring, discovering)
- very large** for a given underlying hardware and network infrastructure, volumes that require special strategies to enable scale-out in storage with no scale-down in processing are **very large**
- data** basically, **facts** describing an **entity** (e.g., the employee with id='123' has name='Jane', the employee with id='456' has name='John'), grouped in **collections** (e.g., Employee) and **related** to one another (e.g., Jane manages John)

Projecting a Coherent Abstract Model


What does it mean?

- A **data model** comprises a set of abstract, domain-independent concepts with which data in the database can be described.
- Such concepts (e.g., primary/foreign keys) induce **integrity constraints** (e.g., referential ones, i.e., a foreign key in one relation must appear as primary key in another).
- A **schema** describes the domain-specific concepts that go into a database in terms of the domain-independent concepts available in the given data model.
- A database **instance** (i.e., the database state at a point in the life cycle of the database) is a snapshot of the world as captured in data and must always be valid with respect to the schema.
- Each DBMS supports one data model (e.g., the relational model), but many supported data models subsume others (e.g., the object-relational model).

High-Level

What does it mean?

- **physical schema**: storage-level structures
- **conceptual/logical schema**: domain-specific, application-independent concepts organized as a collection of **entities** (e.g., relations) and **relationships** (e.g., expressed by means of primary/foreign keys)
- **external schema**: application-specific concepts/requirements as a collection of **views/queries** over the logical schema



Linguistic Abstractions (1)

What does it mean?

A DBMS typically supports three sub-languages:

DDL A **data definition language** to formulate schema-level concepts.

DML A **data manipulation language** to formulate changes to be effected in a database instance.

(D)QL A **(data) query language** to formulate retrieval requests over a database instance.

Linguistic Abstractions (2)

What does it mean?

The best known DBMS language is SQL, different parts of which serve as DDL, DML and QL for (primarily) relational DBMSs.

- The goal of a DDL is to describe application-independent notions (i.e., at the physical and logical levels).
- The goal of a DML and a (D)QL is describe application-specific notions (i.e., at the view level).

```
● CREATE TABLE
    Employee
    (id CHAR(3),name VARCHAR(30),branch VARCHAR(10))

● INSERT INTO
    Employee
    VALUES ('123','Jane','Manchester')
INSERT INTO
    Employee
    VALUES ('456','John','Edinburgh')

● SELECT id
    FROM   Employee
    WHERE  branch = 'Manchester'

● CREATE VIEW
    ManchesterEmployees
    AS
    SELECT id
    FROM   Employee
    WHERE  branch = 'Manchester'
```

Section 2

DBMS: Languages



Database Languages (1)

Are they any different?

- Database languages are different, by design, from general-purpose programming languages.
- Each sub-language caters for distinct needs that in a general-purpose programming language are not factored out.
 - DDL statements update the stored metadata schema-level concepts and correspond to **variable and function declarations** (but have side-effects).
 - DML statements update the database, i.e., change (equivalently, cause a transition from) one database instance into another, and correspond to **assignments**.
 - DQL expressions do not have side-effects, and correspond to **(pure) expressions**.
- Database languages are designed to operate on collections, without explicit iteration, whereas most general-purpose programming language are founded on explicit iteration.

Database Languages (2)

Are they limited in any way?

- QLs are limited, by design, compared with general-purpose programming languages.
- They are not Turing-complete.
- Even (ANSI) SQL was not Turing-complete until procedural constructs (i.e., so called persistent stored modules) were introduced to make it so.
- They are not meant for complex calculations nor for operating on anything other than large collections.
- These limitations make QLs declarative and give them simple formal semantics (in calculus and in algebraic form).
- These properties in turn make it possible for a declarative query to be rewritten by an optimizer into an efficient procedural execution plan.

Section 3


DBMS-Centred Applications



The Architecture of DBMS-Centred Applications

Three Tiers

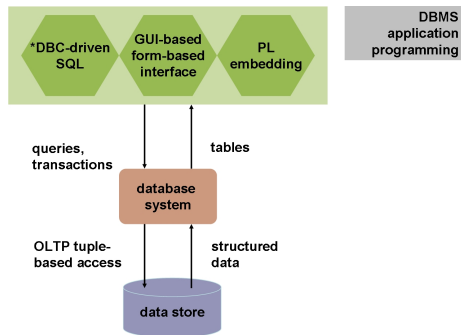
- A classical DBMS supports **on-line transaction processing (OLTP)** applications.
- Applications send queries and transactions that the DBMS converts into OLTP tuple-based operations over the data store.
- What comes back is structured data (e.g., records) as tables, i.e., collections of tuples.



Application Interfaces

Three Routes

- Occasional, non-knowledgeable users are served by graphical user interfaces (GUI), which are often form-based.
- Most application programming tends to benefit from database connectivity middleware (e.g., ODBC, JDBC, JDO, etc.).
- If necessary, from a general-purpose programming language, applications can invoke DBMS services.



Summary

Database Management Systems



- The added-value that DBMSs deliver to organizations stems from controlled use of abstraction.
- The adoption of application-independent data models provides a common formal framework upon which several levels of description become available.
- Such descriptions are conveyed in formal languages that separate concerns and facilitate the implementation of efficient evaluation engines.
- DBMS support OLTP efficiently and do so while scaling along several dimensions.
- They offer convenient interfaces to different classes of users balancing the widening of coverage with the narrowing of needs.

Lecture 2

Database Management Systems [2]



Section 1

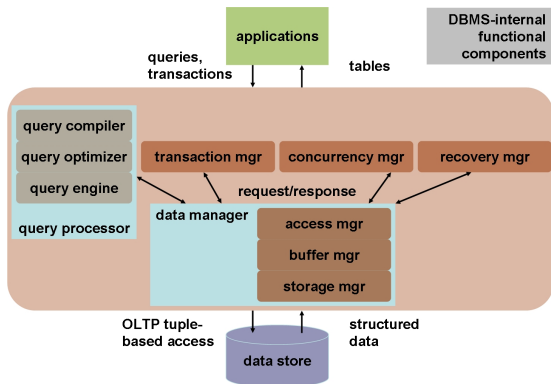
DBMS: Internal Architecture



Internal Architecture of Classical DBMSs

Four+One Main Service Types

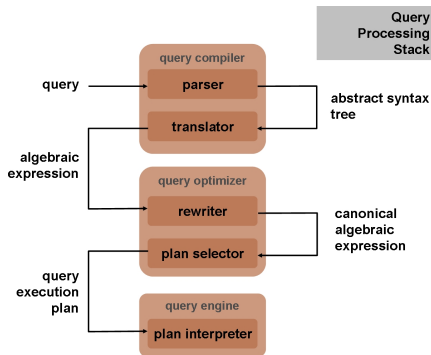
- Query Processing
- Transaction Processing
- Concurrency Control
- Recovery
- Storage Management



The Query Processing Stack

Declarative-to-Procedural, Equivalence-Preserving Program Transformation


- 1 **Parse** the declarative query
- 2 **Translate** to obtain an algebraic expression
- 3 **Rewrite** into a canonical, heuristically-efficient logical query execution plan (QEP)
- 4 **Select** the algorithms and access methods to obtain a quasi-optimal, costed concrete QEP
- 5 **Execute** (the typically interpretable form of) the procedural QEP



Data Store Contents

Four Main Types of Data

- **Metadata** includes schema-level information and a description of the underlying computing infrastructure.
- **Statistics** are mostly information about the trend/summary characteristics of past database instances.
- **Indices** are built for efficient access to the data.
- **Data** is what the database instance contains.



Summary

Classical DBMSs



- Classical DBMSs have been incredibly successful in underpinning the day-to-day life of organizations.
- Their internal functional architecture had not, until very recently, changed much over the last three decades.
- More recently, this architecture has been perceived as being unable to deliver certain kinds of services to business.
- Under pressure from business interests and reacting to opportunities created by new computing infrastructures, classical DBMSs have been transforming themselves into different kinds of advanced DBMSs that this course will explore.

Section 2

DBMS: Strengths, Weaknesses, Trends



Classical DBMSs: Strengths

What are classical database management systems good at?

- Classical database management systems (DBMSs) have been very successful.
- In the last four decades, they have become an indispensable infrastructural component of organizations.
- They play a key role in reliably and efficiently reflecting the transaction-level unfolding of operations in the value-adding chain of an organization.
- Each transaction (e.g., an airline reservation, a credit card payment, an item sold in a checkout) is processed soundly, reliably, and efficiently.
- Effects are propagated throughout the organization.

Classical DBMSs: Weaknesses

Where do classical database management systems come short?

- Classical database management systems (DBMSs) assume that:
 - ① Data is structured in the form of records
 - ② Only on-line transaction processing (OLTP) is needed
 - ③ Data and computational resources are centralized
 - ④ There is central control over central resources
 - ⑤ There is no need for dynamically responding in real-time to external events
 - ⑥ There is no need for embedding in the physical world in which organizations exist
- This is too constraining for most modern businesses.
- Classical DBMSs support fewer needs of organizations than they used to.
- “The Web changes everything.”

Classical DBMSs: Trends

How are DBMSs evolving?

- Most cutting-edge research in databases is geared towards supporting:
 - ① Un- and semi-structured data too
 - ② On-line analytical processing (OLAP) too
 - ③ Distributed data and computational resources
 - ④ Absence of central control over distributed resources
 - ⑤ Dynamic response in real-time to external events
 - ⑥ Embedding in the physical world in which the organization exists
- DBMSs that exhibit these capabilities are **advanced** in the sense used here.

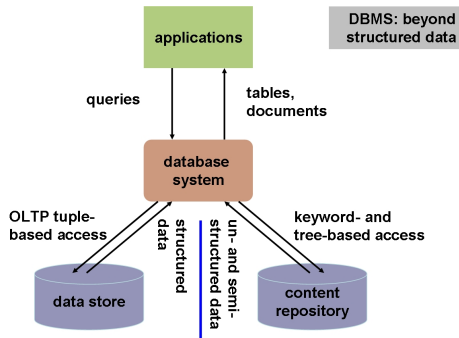
Section 3

DBMS: Architectural Variations



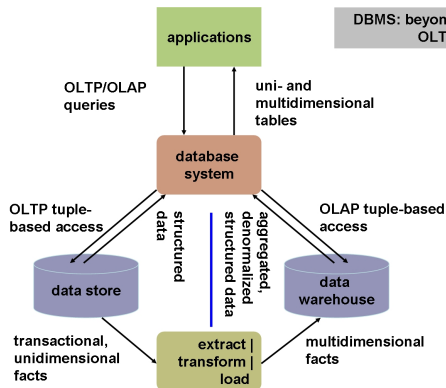
Beyond Structured Data

- Add support to un- and semi-structured data in document form
- Stored in content repositories
- Using keyword-based search and access methods for graph/tree fragments



Beyond OLTP

- Preprocess, aggregate and materialize separately
- Add support for OLAP
- Using multidimensional, denormalized logical schemas




Parallelization (1)

Shared-Disk Parallelism



- Place a fast interconnect between memory and comparatively slow disks
- Parallelize disk usage to avoid secondary-memory contention




Parallelization (2)

Shared-Memory Parallelism



- Place a fast interconnect between processor and memory
- Parallelize memory usage to avoid primary-memory contention

Parallel DBMS:
shared-memory
parallelism




Parallelization (3)

Shared-Nothing Parallelism



- Place a fast interconnect between full processing units
- Parallelize processing using black-boxes that are locally resource-rich




Distribution (1)

Multiple DBMSs



- Harness distributed resources
- Using a full-fledged local or wide-area network as interconnect

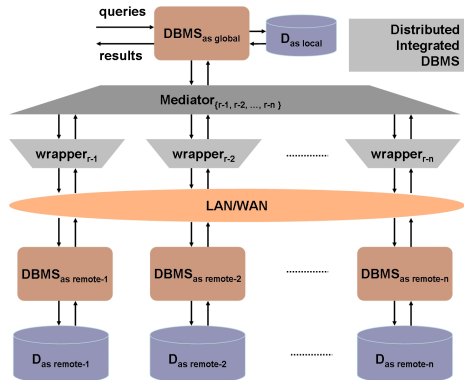


Distribution (2)

Global, Integrated DBMSs



- Renounce central control
- Harness heterogeneous, autonomous, distributed resources
- Project a global view by mediation over wrapper-homogenized local sources




Distribution (3)

Peer-to-Peer DBMSs



- Renounce global view and query expressiveness
- Benefit from inherent scalability over large-scale, extremely-wide-area networks




Distribution (4)

Data Stream Management Systems



- Enable dynamic response in real-time to external events
- Placing queries that execute periodically or reactively
- Over data that is pushed onto the system in the form of unbounded streams




Distribution (5)

Sensor Network Data Management



- Embed data-driven processes in the physical world
- Overlaying query processing over an ad-hoc wireless network of intelligent sensor nodes
- Over pull-based data streams



Advanced DBMSs (1)

Why do they matter?

OLAP/DM Companies need to make more, and more complex, decisions more often and more effectively to remain competitive.

Text-/XML-DBMSs The ubiquity and transparency of networks means data can take many forms, is everywhere, and can be processed anywhere.

Advanced DBMSs (2)

Why do they matter?

P/P2P/DDBMSs For both data and computation, provision of resources is now largely servicized and can be negotiated, or harvested.

Stream DMSs Widespread cross-enterprise integration means that companies must be able to respond in real-time to events streaming in from their commercial and financial environment

Sensor DMSs Most companies are aiming to sense and respond not just to the commercial and financial environment but to the physical environment too.

Summary

Advanced DBMSs



- Architecturally, advanced DBMSs characterize different responses to
 - modern functional and non-functional application requirements
 - the availability of advanced computing and networking infrastructures
- Advanced DBMSs are motivated by real needs of modern organizations, in both the industrial and the scientific arena.

Lecture 3

The Relational Case



Section 1

The Relational Model



The Relational Model of Data

Why?

- Conceptually simple:
 - one single, collection-valued, domain-independent type
- Formally elegant:
 - a very constrained system of first-order logic with both a model- and a proof-theoretic view
 - gives rise to (formally equivalent) declarative and procedural languages, i.e., the domain and the tuple relational calculi, and the relational algebra, resp.
- Practical:
 - underlies SQL
 - possible to implement efficiently
 - has been so implemented many times
- Flexible:
 - often accommodates useful extensions

Section 2

Relational Databases



Relational Databases (1)

Definitions (1)

Definition

A **relational database** is a set of relations.

Example

$D = \{ \text{Students, Enrolled, Courses, } \dots \}$

Definition

A **relation** is defined by its schema and consists of a collection of tuples/rows/records that conform to that schema.

Relational Databases (2)

Definitions (2)

Definition

A **schema** defines the **relation name** and the **name** and **domain**/type of its **attributes**/columns/fields.

Example

Students (stdid: integer, name: string, login: string, age: integer, gpa: real)

Relational Databases (3)

Definitions (3)

Definition

Given its schema, a **relation (instance)** is a subset of the Cartesian product induced by the domain of its attributes.

Definition

A **tuple** in a relation instance is an element in the Cartesian product defined by the relation schema that the instance conforms to.

Example

stdid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Relational Databases (4)

Underlying Assumptions

- Relations are classically considered to be a set (hence, all tuples are unique and their order does not determine identity).
- In practice (e.g., in SQL), relations are multisets/bags, i.e., they may contain duplicate tuples, but their order still does not determine identity.

Relational Databases (5)

Definitions (4)

Definition

The number of attributes in a relation schema defines its **arity**/degree.

Definition

The number of tuples in a relation defines its **cardinality**.

Example

- $\text{arity}(\text{Students}) = 5$
- $\text{cardinality}(\text{Students}) = |\text{Students}| = 3$

Relational Databases (6)

Integrity Constraints (1)

Definition

An **integrity constraint (IC)** is a property that must be true for all database instances.

Example

Domain Constraint: The value of an attribute belongs to the schema-specified domain.

- ICs are specified when the schema is defined.
- ICs are checked when a relation is modified.
- A legal instance of a relation is one that satisfies all specified ICs.
- A DBMS does not normally allow an illegal instance to be stored (or to result from an update operation).

Relational Databases (7)

Integrity Constraints (2)

Definition

- ① A set of fields is a **key** for a relation if both:
 - ① No two distinct tuples can have the same values for those fields.
 - ② This is not true for any subset of those fields.
- ② A **superkey** is not a key, it is a set of fields that properly contains a key.
- ③ If there is more than one key for a relation, each such key is called a **candidate key**.
- ④ The **primary key** is uniquely chosen from the candidate keys.

Relational Databases (8)

Integrity Constraints (3)

Example

- $\{\text{stdid}\}$ is a key in Students, so is $\{\text{login}\}$, $\{\text{name}\}$ is not.
- $\{\text{stdid}, \text{name}\}$ is a superkey in Students.
- $\{\text{stdid}\}$ and $\{\text{login}\}$ are candidate keys in Students
- $\{\text{stdid}\}$ may have been chosen to be the primary key out of the candidate keys.

Relational Databases (9)

Integrity Constraints (4)

- A **foreign key** is set of fields in one relation that
 - appears as the primary key in another relation
 - can be used to refer to tuples in that other relation
 - by acting like a logical pointer
 - it expresses a relationship between two entities
- A DBMS does not normally allow an operation whose outcome violates referential integrity.

Relational Databases (10)

Integrity Constraints (5)

Example

Enrolled (cid: string, grade: string, studentid: string)

Example

cid	grade	studentid
CS101	A	53666
MA102	B	53688
BM222	B	53650

- E.g. {studentid} is a foreign key in Enrolled using the {stdid} primary key of Students to refer to the latter.

Relational Databases (11)

Where Do Integrity Constraints Come From?



- ICs are an aspect of how an organization chooses to model its data requirements in the form of database relations.
- Just like a schema must be asserted, so must ICs.
- We can check a database instance to see if an IC is violated, but we cannot infer that an IC is true by looking at an instance.
- An IC is a statement about all possible instances, not about the particular instance we happen to be looking at.

Summary

Relational Databases



- Since their conception in the late 1960s, and particularly after the first successful, industrial-strength implementations appeared in the 1970s, relational databases have attracted a great deal of praise for their useful elegance.
- Over the 1980s, relational databases became the dominant paradigm, a position they still hold (with some notable evolutionary additions).

Lecture 4

Relational Query Languages [1]



Section 1

Query Languages: Special Features



Relational Query Languages (1)

Why do they matter?

- They were a novel contribution and are a major strength of the relational model.
- They support simple, powerful, well-founded querying of data.
- Requests are specified declaratively and delegated to the DBMS for efficient evaluation.

Relational Query Languages (2)

Why are they special?

- The success of this approach depends on
 - the definition of a pair of query languages, one declarative and one procedural
 - being given a formal semantics that
 - allows their equivalence to be proved
 - the mapping from one to other to be formalized
 - with closure properties (i.e., any expression evaluates to an output of the same type as its arguments) for recursive composition.
- In view of such results, the DBMS can implement a domain-independent query processing stack, such as we have seen in a previous lecture.

Relational Query Languages (3)

Declarative and Procedural, Abstract and Concrete

- By **declarative** we mean a language in which we describe the desired answer without describing how to compute it.
- By **procedural** we mean a language in which we describe the desired answer by describing how to compute it.
- By **abstract** we mean that the language does not have a concrete (let alone, standardized) syntax (and thus, no reference implementation).
- By **concrete** we mean that the language does have a concrete (ideally, standardized) syntax (and thus, often a reference implementation as well).

Relational Query Languages (4)

Domain/Tuple Relational Calculi, Relational Algebra, SQL

- Classically, the relational model defines three expressively-equivalent abstract languages:
 - the **domain relational calculus (DRC)**
 - the **tuple relational calculus (TRC)**
 - the **relational algebra (RA)**
- RA is procedural, DRC and TRC are declarative.
- SQL** (for **Structured Query Language**) is a concrete language whose core is closely related to TRC.

Section 2

Relational Languages Compared, and Views



Relational Query Languages (5)

A First Glimpse

Example

Retrieve the gpa of students with age greater than 18.

TRC: $\{A \mid \exists S \in Students (S.age > 18 \wedge A.gpa = S.gpa)\}$

RA: $\pi_{gpa}(\sigma_{age>18}(Students))$

SQL: `SELECT S.gpa FROM Students S WHERE S.age > 18`

Answer:

gpa

3.8

Relational Query Languages (6)

What do these queries compute?

Example

TRC: $\{A \mid \exists S \in Students \exists E \in Enrolled (S.stdid = E.studentid \wedge E.grade \neq 'B' \wedge A.name = S.name \wedge A.cid = E.cid)\}$

RA: $\pi_{name, cid}(\sigma_{grade \neq 'B'}(Students \bowtie_{stdid=studentid} Enrolled))$

SQL: SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.stdid = E.studentid AND E.grade <> 'B'

Retrieve the names of the students who had a grade different from 'B' and the course in which they did so. Answer:

name	cid
Jones	CS101

Relational Query Languages (7)

Views as Named Queries



- A relation instance is normally defined extensionally (i.e., at each point in time we can enumerate the tuples that belong to it).
- A view defines a relation instance intensionally (i.e., by means of an expression that, when evaluated against a database instance, produces the corresponding relation instance).
- A view explicitly assigns a name to the relation it defines and implicitly characterizes its schema (given that the type of the expression can be inferred).
- Typically, (the substantive part of) the view definition language is the (D)QL,
- For a view, therefore, the DBMS only need store the query expression, not a set of tuples, as the latter can be obtained, whenever needed, by evaluating the former.

Relational Query Languages (8)

View Definition and Usage



- Start with:

Example

```
CREATE VIEW TopStudents (sname, stid, courseid)
  AS SELECT S.name, S.stdid, E.cid
     FROM   Students S, Enrolled E
     WHERE  S.stdid = E.studentid and E.grade = 'B'
```

- Follow up with:

Example

```
SELECT T.sname, T.courseid
FROM   TopStudents T
```

- This should look familiar.

Relational Query Languages (9)

Why are views useful?



- Views can be used to present necessary information (or a summary thereof), while hiding details in underlying relation(s).
- Views can be used to project specific abstractions to specific applications.
- Views can be materialized (e.g., in a data warehouse, to make OLAP feasible).
- By 'materializing' a view we mean evaluating the view and storing the result.
- Views are a useful mechanism for controlled fragmentation and integration in distributed environments.

Summary

Relational Model, Databases, Query Languages



- The relational model remains the best formal foundation for the study of DBMSs.
- It brings out the crucial role of query languages in providing convenient mechanisms for interacting with the data.
- It lies behind the most successful DBMSs used by organizations today.
- For a clear and comprehensive formal description of the relational query languages, see Ch. 6 of [Silberschatz et al., 2009].

Section 3

The Tuple Relational Calculus



General Form and Basic Notation

- A query in the tuple relational calculus is of the form

$$\{t|P(t)\}$$

- It denotes the set of all tuples t such that the predicate P is true of t .
- To denote a value of tuple t on attribute A , we write $t.A$ (but the notation $T[A]$ is also common), often referred to as a tuple access expression.
- We write $t \in R$ to denote that t is a tuple in relation R .

An Example Revisited

Example

TRC: $\{A \mid \exists S \in Students \exists E \in Enrolled (S.stdid = E.studentid \wedge E.grade \neq 'B' \wedge A.name = S.name \wedge A.cid = E.cid)\}$

ENG: *Retrieve the names of the students who had a grade different from 'B' and the course in which they did so.*

- Note how the idiom $t \in R$ is used to denote an iteration, and how more than one such usage denotes nested iteration.
- Note how the predicate has two logical parts:
 - In the first part, i.e., the first two conjuncts, we express the conditions that result tuples must satisfy. Note that they do not have tuple access expressions on the left-hand side.
 - In the second part, i.e., the last two conjuncts, we express **variable bindings**, using tuple access expressions on the left-hand side to define which values each attribute in a result tuple will have.

The Procedural Meaning of a TRC Expression (1)

Data on Beers, Bars and Drinkers

Example

```
Beer = {0: ('Name','City'),  
        1: ('peer', 'edi'),  
        2: ('pier', 'man'),  
        3: ('pear', 'man'),  
        4: ('pare', 'man')}
```

```
cardBeer = len(Beer)-1
```

```
Bar = { 0: ('Name','City'),  
        1: ('boar', 'lon'),  
        2: ('bear', 'edi'),  
        3: ('bare', 'man'),  
        4: ('bore', 'man')}
```

```
cardBar = len(Bar)-1
```

```
Drinker = { 0: ('Name','City'),  
            1: ('jane','man'),  
            2: ('jon', 'lon')}
```

```
cardDrinker = len(Drinker)-1
```

The Procedural Meaning of a TRC Expression (2)

Expressing the query: Retrieve cities that have both bars and drinkers

Example

```
Q = """
{R | exists d in drinker
    exists b in bar
    (d.City = b.City and
     R.BarAndDrinker = d.City
    )
}
"""
```

The Procedural Meaning of a TRC Expression (3)

Computing the Result

Example

```
# initialize the result relation R and set its schema
R = {0: ('BarAndDrinker')}

# initialize the tuple id counter for R
k = 0

# one FOR loop per existential quantifier
for i in range(1,cardDrinker+1):
    for j in range(1,cardBar+1):
        # fetch the values in the scope of the quantifiers
        d = Drinker[i]
        b = Bar[j]
        # write out the condition part of the conjunction
        if d[1] == b[1]:
            # set semantics requires duplicate elimination
            if not (d[1] in R.values()):
                # new tuple id
                k = k+1
                # write out the binding part of the conjunction
                # using tuple accessor expressions
                R[k] = d[1]

# return the result
return R
```

Section 4

A Relational Algebra [1]



Preliminaries

- A query is applied to relation instances.
- The result of a query is also a relation instance.
- The schemas of the input/argument relations of a query are fixed.
- The schema for the result of a given query is statically known by type inference on the schemas of the input/argument relations.
- Recall that relational algebra is **closed** (equiv., has a closure property), i.e., the input(s) and output of any relational-algebraic expression is a relation instance.

Relational Algebra (1)

Relation Instances Used in Examples (1)



- Let Sailors, Boats and Reservations be example relations.
- Their schemas are on the right.
- Underlined sets of fields denote a key.

Example

Sailors (sid: integer, sname: string, rating: integer, age: real)

Boats (bid: integer, bname: string, colour: string)

Reservations (sid: integer, bid: integer, day: date)

Relational Algebra (2)

Relation Instances Used in Examples (2)



- The various relation instances used on the right.
- Note that there are two relation instances (viz., S1 and S2) for Sailors, one for Reservations (viz., R1), and none, yet, for Boats.
- Fields in an instance of one of these relations are referred to by name or by position (in which case the order is that of appearance, left to right, in the

Example

S1 =

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2 =

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

R1 =

sid	bid	day
22	101	10/10/96
58	103	12/11/96

Relational Algebra (3)

Primitive Operations

- σ : **selection** returns those rows in the single argument-relation that satisfy the given predicate
- π : **projection** deletes those columns in the single argument-relation that are not explicitly asked for
- \times : **Cartesian (or cross) product** concatenates each row in the first argument relation with each row in the second to form a row in the output
- \setminus : **(set) difference** returns the rows in the first argument relation that are not in the second
- \cup : **(set) union** returns the rows that are either in the first or in the second argument relation (or in both)

The above is a complete set: any other relational-algebra can be derived by a combination of the above.

Relational Algebra (4)

Derived Operations

$$\cap : R \cap S \Leftrightarrow (R \cup S) \setminus ((R \setminus S) \cup (S \setminus R))$$

(set) intersection returns the rows that are both in the first and in the second argument relation

$$\bowtie : R \bowtie_{\theta} S \Leftrightarrow \sigma_{\theta}(R \times S)$$

join concatenates each row in the first argument relation with each row in the second and forms with them a row in the output, provided that it satisfies the given predicate

$$\div : (\text{see below for derivation})$$

division returns the rows in the first argument relation that are associated with every row in the second argument relation

Relational Algebra (5)

Extensions (1)

Useful extensions for clarity of exposition are:

- ρ : **renaming** returns the same relation instance passed as argument but assigns the given name(s) to the output relation (or any of its attributes)
- \leftarrow : **assignment** assigns the left-hand side name to the relation instance resulting from evaluating the right-hand side expression

Relational Algebra (6)

Extensions (2)

Extensions that change the expressiveness of classical relation algebra include (see, e.g., [Silberschatz et al., 2009]):

- **generalized projection**, which allows arithmetic expressions (and not just attribute names) to be specified
- **(group-by) aggregation**, which allows functions (such as **count**, **sum**, **max**, **min**, **avg**) to be applied on some attribute (possibly over partitions defined by the given group-by attribute)
- other kinds of join (e.g., semijoin, antijoin, [left|right] outer join)

Lecture 5

Relational Query Languages [2]



Section 1

A Relational Algebra [2]



Selection

- $\sigma_{\theta}(R) = \{x \mid \exists x \in R (\theta(x))\}$
- Rows in the single input relation R that satisfy the given selection condition (i.e., a Boolean expression on the available attributes) are in the output relation O .
- No duplicate rows can appear in O , so the cardinality of O cannot be larger than that of R .
- The schema of O is identical to the schema of R .
- The arity of O is the same as that of R .

Example

$$\sigma_{rating > 8}(S2) =$$

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{sid > 10 \wedge age < 45.0}(\sigma_{rating > 8}(S2)) =$$

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

Projection

- $\pi_{a_1, \dots, a_n}(R) = \{y \mid \exists x \in R (y.a_1 = x.a_1 \wedge \dots \wedge y.a_n = x.a_n)\}$
- Columns in the single input relation R that are not in the given projection list do not appear in the output relation O .
- Duplicate rows might appear in O unless they are explicitly removed, and, if so, the cardinality of O is the same as that of R .
- The schema of O maps one-to-one to the given projection list, so the arity of O cannot be larger than that of R .

Example

$\pi_{sname, rating}(S2) =$

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(\sigma_{rating > 8}(S2)) =$

sname	rating
yuppy	9
rusty	10

Set Operations (1)

Union, Intersection, Difference



- These binary operations have the expected set-theoretic semantics.
- Both input arguments R and S must have compatible schemas (i.e., their arity must be the same and the columns have to have the same types one-to-one, left to right).
- The arity of O is identical to that of I .
- The cardinality of O may be larger than that of the largest between R and S in the case of union, but not in the case of intersection and difference.

Example

$S1 \cup S2 =$

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
28	yuppy	9	35.0
44	guppy	5	35.0

Set Operations (2)

Union, Intersection, Difference



Example

$S1 \cap S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

Example

$S1 \setminus S2$

sid	sname	rating	age
22	dustin	7	45.0

Cartesian/Cross Product

- $R \times S = \{xy \mid \exists x \in R \exists y \in S\}$
- The schema of O is the concatenation of the schemas of R and S , unless there is a name clash, in which case renaming can be used.
- The arity of O is the sum of the arities of R and S .
- The cardinality of O is the product of the cardinalities of R and S .

Example

$\rho_{1 \rightarrow sid1, 5 \rightarrow sid2}(S1 \times R1) =$

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	12/11/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	12/11/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	12/11/96

Joins (1)

θ -Join

- $R \bowtie_{\theta} S = \{xy \mid \exists x \in R \exists y \in S (\theta(xy))\}$
- $R \bowtie_{\theta} S \equiv \sigma_{\theta}(R \times S)$
- The schema of O is as for Cartesian product, as is arity.
- The cardinality of O cannot be larger than the product of the cardinalities of R and S .

Example

$\rho_{1 \rightarrow \text{sid1}, 5 \rightarrow \text{sid2}}(S1 \bowtie_{S1.\text{sid} < R1.\text{sid}} R1) =$

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	58	103	12/11/96
31	lubber	8	55.5	58	103	12/11/96

Joins (2)

Equijoin and Natural Join

- An **equijoin** is a θ -join in which all terms in θ are equalities.
- In an equijoin, the schema of O is as for Cartesian product but only one of the equated columns is projected out, so the arity reduces by one for each such case.
- A **natural join** is an **equijoin** on all common columns.
- One only needs list the common columns in the condition.

Example

$S1 \bowtie_{sid} R1 =$

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	12/11/96

Division (1)

Through Examples



- In integer division, given integers A and B , $A \div B$ is the largest integer Q such that $Q \times B \leq A$.
- In relational division, given two relations R and S , $R \div S$ is the largest relation instance O such that $O \times S \subseteq R$.
- If R lists suppliers and parts they supply, and S parts, then $R \div S$ lists suppliers of all S parts.

Example

$A =$

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

Example

$B_1 =$

pno
p2

$B_2 =$

pno
p2
p4

$B_3 =$

pno
p1
p2
p4

Example

$A \div B_1 =$

sno
s1
s2
s3
s4

$A \div B_2 =$

sno
s1
s4

$A \div B_3 =$

sno
s1

Division (2)

Through Rewriting (1)



- Division, like join, can be defined by rewriting into primitive operations but, unlike join, it is not used very often, so most DBMSs do not implement special algorithms for it.
- The schema of O is the schema of R minus the columns shared with S , so the arity of O cannot be as large as that of R .
- The cardinality of O cannot be larger than that of R .

Division (3)

Through Rewriting (2)



- Abusing notation, we can define division in terms of primitive operators as follows.
- Let r and s be relations with schemas R and S , respectively, and let $S \subseteq R$, then:
 - $T_1 \leftarrow \pi_{R-S}(r) \times s$ computes the Cartesian product of $\pi_{R-S}(r)$ and s so that each tuple $t \in \pi_{R-S}(r)$ is paired with every s -tuple.
 - $T_2 \leftarrow \pi_{R-S,S}(r)$ merely reorders the attributes of r in preparation for the set operation to come.
 - $T_3 \leftarrow \pi_{R-S}(T_1 - T_2)$ only retains those tuples $t \in \pi_{R-S}(r)$ such that for some tuple u in s , $tu \notin r$.
 - $r \div s \leftarrow \pi_{R-S}(r) - T_3$ only retains those tuples $t \in \pi_{R-S}(r)$ such that for all tuples u in s , $tu \in r$.

Generalized Projection

- Recall that generalized projection allows arithmetic expressions involving attribute names (and not just the latter) in the projection list.
- The first example to the right returns the sailor names with their associated ranking tripled.
- There is a further extended version that allows concomitant renaming as shown in the second example to the right.

Example

$$\pi_{sname, rating * 3}(S2) =$$

sname	rating
yuppy	27
lubber	24
guppy	15
rusty	30

$$\pi_{sname, rating * 3 \rightarrow triplerating}(S2) \equiv$$

$$\rho_{2 \rightarrow triplerating}(\pi_{sname, rating * 3}(S2)) =$$

sname	triplerating
yuppy	27
lubber	24
guppy	15
rusty	30

Aggregation

- Recall that aggregation reduces a collection of values into a single values by the application of a function such as **count**, **sum**, **max**, **min** or **avg**.
- It is also possible to form groups by attribute values, e.g., to take the average rating by age.
- Concomitant renaming can also be used.

Example

$\gamma_{avg(age) \rightarrow averageage}(S2) =$

averageage

40.125

$age \gamma_{avg(rating) \rightarrow averagerating}(S2) =$

age	averagerating
35.0	8
55.5	8

Section 2

RA Examples



Example Relational Algebra Expressions (1)

More Relation Instances



- For the next batch of examples, the various relation instances used are as shown.

Example

S3 =

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5

Example

R2 =

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	08/10/98
22	104	07/10/98
31	102	10/11/98
31	103	06/11/98
31	104	12/11/98
64	101	05/09/98
64	102	08/09/98
64	103	08/09/98

B1 =

bid	bname	colour
101	interlake	blue
102	interlake	red
103	clipper	green
104	marine	red

Example Relational Algebra Expressions (2)

Find the names of the sailors who have reserved boat 103



- $O_1 = \pi_{sname}(\sigma_{bid=103}(Reservations \bowtie Sailors))$
- $O_2 = \pi_{sname}((\sigma_{bid=103}(Reservations)) \bowtie Sailors)$
- $O_1 \equiv O_2$

Example

sname
dustin
lubber
horatio

Example Relational Algebra Expressions (3)

Find the names of the sailors who have reserved a red boat



- Information about boat colour is only available in Boats, so an extra join is needed.
- $O_1 = \pi_{sname}(\sigma_{color=red}(Boats) \bowtie (Reservations \bowtie Sailors))$
- $O_2 = \pi_{sname}(\pi_{sid}(\pi_{bid}(\sigma_{color=red}(Boats)) \bowtie Reservations) \bowtie Sailors)$
- $O_1 \equiv O_2$

Example

sname
dustin
lubber
horatio

Example Relational Algebra Expressions (4)

Find the names of the sailors who have reserved a red or a green boat



- Using assignment:

$$\textcircled{1} \quad T_1 \leftarrow (\sigma_{\text{colour}=\text{red}}(\text{Boats}) \cup \sigma_{\text{colour}=\text{green}}(\text{Boats}))$$

$$\textcircled{2} \quad O \leftarrow \pi_{\text{sname}}(T_1 \bowtie (\text{Reservations} \bowtie \text{Sailors}))$$

- Or:

$$\textcircled{1} \quad T_1 \leftarrow (\sigma_{\text{colour}=\text{red} \vee \text{colour}=\text{green}}(\text{Boats}))$$

$$\textcircled{2} \quad O \leftarrow \pi_{\text{sname}}(T_1 \bowtie (\text{Reservations} \bowtie \text{Sailors}))$$

Example

sname

dustin

lubber

horatio

Example Relational Algebra Expressions (5)

Find the names of the sailors who have reserved a red and a green boat



- Replacing \cup with \cap in the previous example doesn't work.
- Using assignment:
 - ① $T_1 \leftarrow \pi_{sid}(\sigma_{colour=red}(Boats) \bowtie Reservations)$
 - ② $T_2 \leftarrow \pi_{sid}(\sigma_{colour=green}(Boats) \bowtie Reservations)$
 - ③ $O \leftarrow \pi_{sname}((T_1 \cap T_2) \bowtie Sailors)$
- On the other hand, $\pi_{sname}((T_1 \cup T_2) \bowtie Sailors)$ does work for the previous example.

Example

sname
dustin
lubber
horatio

Example Relational Algebra Expressions (6)

Find the ids of sailors older than 20 who have not reserved a red boat



- Using assignment:

- 1 $T_1 \leftarrow \pi_{sid}(\sigma_{age>20}(Sailors))$
- 2 $T_2 \leftarrow \pi_{sid}(\sigma_{colour=red}(Boats) \bowtie Reservations)$
- 3 $O \leftarrow T_1 \setminus T_2$

Example

sid
29
32
58
74
85
95

Example Relational Algebra Expressions (7)

Find the names of sailors who have reserved all boats



- The word **all** suggests the need for division.
- Projections are essential to arrange the schemas appropriately.
- Joins may be needed to recover columns that had to be dropped.
- - 1 $T_1 \leftarrow \pi_{sid,bid}(Reservations) \div \pi_{bid}(Boats)$
 - 2 $O \leftarrow \pi_{sname}(T_1 \bowtie Sailors)$

Example

sname

dustin

Summary

Relational Algebra



- An algebra, often extending, or modelled on, the relational algebra lies at the heart of most advanced DBMSs.
- It is the most used target formalism for the internal representation of logical plans.
- Rewriting that is based on logical-algebraic equivalences is an important task in query optimization (as we will discuss later).

Section 3

SQL



SQL

Relation Instances Used in Examples



- Recall the relation instances on the right.
- These have been used before and will be used again, as before, in the examples that follow.

Example

S1 =

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2 =

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

R1 =

sid	bid	day
22	101	10/10/96
58	103	12/11/96

Core, Informal SQL Syntax (1)

The SELECT, FROM and WHERE Clauses

Definition

```
SELECT  [DISTINCT] < target list >  
FROM    < relation list >  
WHERE   < qualification >
```

- The SELECT clause defines which columns participate in the result, i.e., it plays the role of the relational-algebraic π operation.
- The FROM clause defines which relations are used as inputs, i.e., it corresponds to the leaves in a relational-algebraic expression.
- The WHERE clause defines the (possibly complex) predicate expression which a row must satisfy to participate in the result, i.e., it supplies the predicates for relational-algebraic operations like σ and \bowtie .
- DISTINCT is an optional keyword indicating that duplicates must be removed from the answer.

Core, Informal SQL Syntax (2)

The Arguments to SELECT, FROM and WHERE Clauses

- The argument to a FROM clause is a list of relation names (possibly with a range variable after each name, which allows a row in that relation to be referred to elsewhere in the query).
- It is good practice to use range variables, so use them.
- The argument to a SELECT clause is a list of expressions based on attributes taken from the relations in the **relation list**.
- If '*' is used instead of a list, all available attributes are projected out.
- The argument to a WHERE clause is referred to as a **qualification**, i.e., a Boolean expression whose terms are comparisons (of the form $E \text{ op } \text{const}$ or $E_1 \text{ op } E_2$ where E, E_1, E_2 are, typically, attributes taken from the relations in the **relation list**, and $\text{op} \in \{<, >, =, >=, =<, <>\}$) combined using the connectives AND, OR and NOT.

Core, Informal SQL Semantics (1)

Three-to-Four Steps to the Answer

- To characterize the semantics of a SQL query using a direct, clause-by-clause translation into a relational-algebraic expression that evaluates to the correct answer, do the following:
 - ① Compute the cross-product of relations in the FROM list, call it J.
 - ② Discard tuples in J that fail the qualification, call the result S.
 - ③ Delete from S any attribute that is not in the SELECT list, call the result P.
 - ④ If DISTINCT is specified, eliminate duplicate rows in P to obtain the result A, otherwise $A=P$.
- While as an evaluation strategy, the procedure above is likely to be very inefficient, it provides a simple, clear characterization of the answer to a query.
- As we will see, an optimizer is likely to find more efficient evaluation strategies to compute the same answer.

Core, Informal SQL Semantics (2)

Find the names of the sailors who have reserved boat 103.

Example

```
SELECT S.sname
FROM Sailors S, Reservations R
WHERE S.sid = R.sid AND R.bid = 103
```

Example

Assume the database state contains $\{S1, R1\}$. Then, in Step 1,
 $S1 \times R1 =$

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	12/11/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	12/11/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	12/11/96

In Step 2, $\sigma_{S.sid=R.sid \wedge R.bid=103}(S1 \times R1)$

sid1	sname	rating	age	sid2	bid	day
58	rusty	10	35.0	58	103	12/11/96

In Step 3, $\pi_{sname}(\sigma_{S.sid=R.sid \wedge R.bid=103}(S1 \times R1))$

sname
rusty

No DISTINCT implies no Step 4.

Core, Informal SQL Semantics (3)

Find the ids of sailors who have reserved at least one boat

Example

```
SELECT S.sid
FROM Sailors S, Reservations R
WHERE S.sid = R.sid
```

Example

Assume the database state contains $\{S1, R1\}$. Then, Step 1, $S1 \times R1$ produces the same results as in the previous example.

In Step 2, $\sigma_{S.sid=R.sid}(S1 \times R1)$

sid1	sname	rating	age	sid2	bid	day
22	dustin	7	45.0	22	101	10/10/96
58	rusty	10	35.0	58	103	12/11/96

In Step 3, $\pi_{sid}(\sigma_{S.sid=R.sid}(S1 \times R1))$

sid
22
58

No DISTINCT implies no Step 4.

Core, Informal SQL Syntax (3)

Expressions and Strings



- Arithmetic expressions and string pattern matching can also be used.
- AS and = are two ways to name fields in result.
- LIKE is used for string matching.
'_' stands for any one character and '%' stands for zero or more arbitrary characters.

Example

```
SELECT  S.age,  
        age1=S.age-5,  
        2*S.age AS age2  
FROM    Sailors S  
WHERE   S.sname LIKE 'y_%y'
```

Over S2, the answer would be:

age	age1	age 2
35.0	30.0	70.0

Example SQL Queries (1)

More Relation Instances



- For the next batch of examples, the various relation instances used are also known from previous examples.

Example

S3 =

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5

Example

R2 =

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	08/10/98
22	104	07/10/98
31	102	10/11/98
31	103	06/11/98
31	104	12/11/98
64	101	05/09/98
64	102	08/09/98
64	103	08/09/98

B1 =

bid	bname	colour
101	interlake	blue
102	interlake	red
103	clipper	green
104	marine	red

Example SQL Queries (2)

Find the names of the sailors who have reserved a red boat



- Recall $\pi_{sname}(\sigma_{color=red}(Boats) \bowtie (Reservations \bowtie Sailors))$

Example

```
SELECT  S.sname
FROM    Sailors S, Boats B, Reservations R
WHERE   S.sid = R.sid
        AND R.bid = B.bid
        AND B.colour = 'red'
```


Example SQL Queries (3)

Find the sids of the sailors who have reserved a red or a green boat



Example

```

Either:  SELECT  S.sid
         FROM    Sailors S, Boats B, Reservations R
         WHERE   S.sid = R.sid
               AND R.bid = B.bid
               AND (B.colour = 'red' OR B.colour = 'green')

Or:      SELECT  S1.sid
         FROM    Sailors S1, Boats B1, Reservations R1
         WHERE   S1.sid = R1.sid
               AND R1.bid = B1.bid
               AND B1.colour = 'red'

         UNION
         SELECT  S2.sid
         FROM    Sailors S2, Boats B2, Reservations R2
         WHERE   S2.sid = R2.sid
               AND R2.bid = B2.bid
               AND B2.colour = 'green'
  
```

Set difference is captured by EXCEPT, e.g., to find the sids of the sailors who have reserved a red but not a green boat.

Example SQL Queries (4)

Beware being quick when there are quirks



Example

Contrast this query:	<pre> SELECT S.sid FROM Sailors S, Boats B, Reservations R WHERE S.sid = R.sid AND R.bid = B.bid AND (B.colour = 'red' AND B.colour = 'green') </pre>
with this one:	<pre> SELECT S.sid FROM Sailors S, Boats B, Reservations R WHERE S.sid = R.sid AND R.bid = B.bid AND B.colour = 'red' INTERSECT SELECT S.sid FROM Sailors S, Boats B, Reservations R WHERE S.sid = R.sid AND R.bid = B.bid AND B.colour = 'green' </pre>
and this one:	<pre> SELECT S.sid FROM Sailors S, Boats B1, Reservations R1, Boats B2, Reservations R2, WHERE S.sid = R1.sid AND S.sid=R2.Sid AND R1.bid = B1.bid AND R2.Bid = B2.Bid AND (B1.colour = 'red' AND B2.colour = 'green') </pre>

Core, Informal SQL Syntax (4)

Nested Queries



Example

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT  R.sid
                   FROM    Reservations R
                   WHERE   R.bid = 103
                 )
```

- The ability to nest queries is a powerful feature of SQL.
- Queries can be nested in the WHERE, FROM and HAVING clauses.
- To understand the semantics of nested queries, think of a nested loop, i.e., for each Sailors tuple, compute the nested query and which pass the IN qualification.

Core, Informal SQL Syntax (5)

Correlated Nested Queries



- We could have used NOT IN to find sailor who have not reserved boat 103.
- Another set comparison operator (implicitly, with the empty set) is EXISTS, and see the Bibliography for yet more.
- It is also possible to correlate the queries via shared range variables.

Example

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS ( SELECT  *
                  FROM    Reservations R
                  WHERE    R.bid = 103 AND R.sid = S.sid
                )
```

Core, Informal SQL Syntax (6)

Aggregate Operators



Definition

COUNT	([DISTINCT] A)	the number of (unique) values in the A column
SUM	([DISTINCT] A)	the sum of all (unique) values in the A column
AVG	([DISTINCT] A)	the average of all (unique) values in the A column
MAX	(A)	the maximum value in the A column
MIN	(A)	the minimum value in the A column

- Aggregate operators are another significant extension of relational algebra in SQL.
- They take a collection of values as input and return a single value as output.

Example SQL Queries (5)

Aggregation Queries



Example

- ```
SELECT COUNT(*)
FROM Sailors S
```
- ```
SELECT AVG(S.age)  
FROM Sailors S  
WHERE S.rating = 10
```
- ```
SELECT COUNT(DISTINCT S.rating)
FROM Sailors S
WHERE S.name = 'horatio' OR S.name = 'dusting'
```
- ```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating = ( SELECT MAX(S2.rating)  
                   FROM Sailors S2  
                   WHERE S2.sname = 'horatio' )
```

Example SQL Queries (6)

Find the name and the age of the oldest sailor(s)



Example

- ```
SELECT S.sname, MAX(S.age) -- Illegal SQL!
FROM Sailors S
```
- ```
SELECT  S.sname, S.age
FROM    Sailors S
WHERE   S.age = ( SELECT  MAX(S2.age)
                  FROM    Sailors S2 )
```

- The first query is illegal: if a SELECT clause uses an aggregate operation either it must do so for all attributes in the clause or else it must contain a GROUP BY clause.
- The second query, with a nested query, is legal and correct.

Core, Informal SQL Syntax (7)

Partitioned Aggregation



Definition

```

SELECT      [DISTINCT] < target list >
FROM        < relation list >
WHERE       < qualification >
GROUP BY    < grouping list >
HAVING      < grouping qualification >
  
```

- A group is a partition of rows that agree on the values of the attributes in the grouping list.
- We can mix attribute names and applications of aggregate operations in the target list but the attribute names must be a subset of the grouping list, so that each row in the result corresponds to one group.
- The grouping qualification determines whether a row is produced in the answer for a given group.

Core, Informal SQL Semantics (4)

Three More Steps Than Before to the Answer



- To characterize what is the answer to this extended form of an SQL query:
 - ① Compute the cross-product of relations in the FROM list, call it J.
 - ② Discard tuples in J that fail the qualification, call the result S.
 - ③ Delete from S any attribute that is not in the SELECT list, call the result P.
 - ④ Sort P into groups by the value of attributes in the GROUP BY list, call the result G.
 - ⑤ Discard groups in G that fail the grouping qualification, call the result H.
 - ⑥ Generate one answer tuple per qualifying group, call the result T.
 - ⑦ If DISTINCT is specified, eliminate duplicate rows in T to obtain the result A, otherwise $A=T$.

Core, Informal SQL Semantics (5)

Find the age of the youngest sailor that is at least 18, for each rating with at least 2 such sailors



Example

```
SELECT      S.rating, MIN(S.age) AS minage
FROM        Sailors S
WHERE       S.age >= 18
GROUP BY    S.rating
HAVING      COUNT(*) >= 2
```

Assume the database state contains {S3}.

Then, in Step 1, S3 =

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5

Example

After Steps 2 and 3, we have:

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
9	35.0
3	25.5
3	63.5

Core, Informal SQL Semantics (6)

Find the age of the youngest sailor that is at least 18, for each rating with at least 2 such sailors



Example

After Step 4, we have:

rating	age
1	33.0
3	25.5
3	63.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0

Example

After Step 5, we have:

rating	age
3	25.5
3	63.5
7	45.0
7	35.0
8	55.5
8	25.5

After Steps 6 and 7, we have:

rating	minage
3	25.5
7	35.0
8	25.5

Example SQL Queries (7)

For each red boat, find the number of reservations for this boat



Example

```
SELECT    B.bid, COUNT(*) AS reservationCount
FROM      Boats B, Reservations R
WHERE     R.bid = B.bid AND B.colour = 'red'
GROUP BY  B.Bid
```

Null Values



- Field values in a tuple are sometimes unknown (e.g., a rating has not been assigned) or inapplicable (e.g., someone who is not married has no value to go in a 'spouse_name' column).
- SQL provides a special value null for such situations.
- The presence of null complicates many issues.
 - Special operators are needed to check if value is/is not null.
 - Is `rating > 8` true or false when `rating` is equal to null? What about AND, OR and NOT connectives?
 - We need a 3-valued logic (true, false and unknown).
 - The meaning of many constructs must be defined carefully (e.g., WHERE clause eliminates rows that don't evaluate to true).
 - New operators (in particular, outer joins) become possible and are often needed.
- In SQL, null values can be disallowed when columns are defined.

Summary

SQL



- SQL was an important factor in the early acceptance of the relational model because it is more natural than earlier, procedural query languages.
- SQL is relationally complete (in fact, it has significantly more expressive power than relational algebra).
- Even queries that can be expressed in RA can often be expressed more naturally in SQL.
- There are usually many alternative ways to write a query, so an optimizer is needed to find an efficient evaluation plan.
- In practice, users need to be aware of how queries are optimized and evaluated for best results.
- NULL for unknown field values brings many complications

This Week: Summary



- We have revised the basic notions of
 - what a **database management system** (DBMS) is
 - the role that is played by the languages a DBMS supports
 - how DBMS-centred applications are designed
- We have adopted an approach to describing the internal architecture of DBMSs that allowed us to discuss
 - the strengths and weaknesses of classical DBMSs
 - the recent trends in the way organizations use DBMS
 - how architectures have been diversifying recently
- We have revised the various kinds of query languages by looking into
 - the (tuple) relational calculus
 - a relational algebra
 - SQL

Next Week: Topics



- We will look in detail into logical optimization, i.e., heuristic-based rewriting of algebraic plans.
- We will explore some of the basic strategies for designing physical operators and for combining them into query evaluation plans (QEPs).
- We will briefly discuss some of the challenges in cost-based QEP selection, in particular the hard problem of choosing a join order.
- We will find out how data partitioning strategies, allied with the algebraic properties of query languages, make it possible to automatically parallelize QEP execution.

Next week: Learning Objectives



- We will aim to complete our exploration of classical query processing.
- We will aim to be ready to use what we learn on our way to consider querying data on the Web, in the narrower sense, starting with XML/XQuery.

References



Silberschatz, A., Korth, H. F., and Sudarshan, S. (2009).

Database System Concepts.

McGraw-Hill Education - Europe, 6th edition.